

Convolutional Neural Networks for Crack Detection Classification

Using Convolutional Neural Networks for binary classification of concrete with and without cracks

Megan Arnold
DTSA 5511 Final
September 30th, 2025

Problem Overview

- Images contain very useful information, but are difficult to classify as there are potentially many features within a single image.
- Detecting cracks and deviations can enable early detection; thus, enabling timely repairs or replacement before catastrophic and expensive failure
- Challenge: Many features not relevant to cracks are present
- Challenge: Labels are binary for the images, nothing spatial is encode

Cracked Examples:



Uncracked Examples:






Project Goals

- Explore dataset to **understand data diversity** and guide model selections and identifying potential pitfalls and confounding factors
- **Build Baseline CNN Model** based on exploratory data analysis.
- **Evaluate model performance** and attempt to improve architecture and parameters for subsequent models
- **Visualize CNN gradient heatmaps** to determine the strengths and weaknesses of models; thus, guiding future architecture decisions
- Use **Data Augmentation** on the best performing model to determine if that will help the model generalize
- **Evaluate hyperparameter** tuning outcomes on test dataset, via Kaggle Submission
- Make future work **recommendations**

Dataset Overview

- About 7500 cracked and uncracked images
- Balanced dataset between binary classification
- Observations
 - Cracked:
 - Thick and thin cracks
 - Some are hardly visible to the naked eye
 - No spatial labeling
 - Uncracked:
 - Many uncracked images have significant craters


Examples of Data:




Exploratory Data Analysis

- 256 by 256 pixels with 3 color channels
- Balanced Dataset
- Consistent histograms on pixel values.
Reduces the probability of the model learning a shortcut
- Challenges:
 - There appears to be a potential confounding factor of craters in uncracked images
 - *Note: My first baseline model used the craters as an important feature for classification, as observed by the heatmaps

Cracked Histograms:




Uncracked Histograms:



Baseline Convolutional Model

- Rescaling to pixel value range of [0,1]
- 7 Convolutional Layers
 - 32,32,64, 64, 128, 128, 256
- Batch Normalization between odd and even layers
- MaxPooling between even and odd layers
- Global Average Pooling between CNN and classification blocks
- Dense layer of 64 neurons
- Sigmoid activation neuron



Model: "sequential"


| Layer (type) | Output Shape | Param # |
|---|----------------------|---------|
| rescaling (Rescaling) | (None, 256, 256, 3) | 0 |
| conv2d (Conv2D) | (None, 256, 256, 32) | 896 |
| batch_normalization (Batch Normalization) | (None, 256, 256, 32) | 128 |
| conv2d_1 (Conv2D) | (None, 256, 256, 32) | 9248 |
| max_pooling2d (MaxPooling2D) | (None, 128, 128, 32) | 0 |
| conv2d_2 (Conv2D) | (None, 128, 128, 64) | 18496 |
| batch_normalization_1 (Batch Normalization) | (None, 128, 128, 64) | 256 |
| conv2d_3 (Conv2D) | (None, 128, 128, 64) | 36928 |
| max_pooling2d_1 (MaxPooling2D) | (None, 64, 64, 64) | 0 |
| conv2d_4 (Conv2D) | (None, 64, 64, 128) | 73856 |
| batch_normalization_2 (Batch Normalization) | (None, 64, 64, 128) | 512 |
| conv2d_5 (Conv2D) | (None, 64, 64, 128) | 147584 |
| conv2d_6 (Conv2D) | (None, 64, 64, 256) | 295168 |
| global_average_pooling2d (GlobalAveragePooling2D) | (None, 256) | 0 |
| dropout (Dropout) | (None, 256) | 0 |
| dense (Dense) | (None, 64) | 16448 |
| dense_1 (Dense) | (None, 1) | 65 |

=====
Total params: 599585 (2.29 MB)
Trainable params: 599137 (2.29 MB)
Non-trainable params: 448 (1.75 KB)


Baseline Observations

- Test Accuracy: 0.78
- General Observations
 - Started overfitting around epoch 10
 - Significant focus on craters instead of cracks
 - Focus on small local areas instead of global context
 - Some focus on parts of a crack that look like craters
 - No edge artifacts seen


Prob: 0.99997830



Prob: 0.21832037




Prob: 0.99999857



Model and Hyperparameter Tuning Version 2

- Overfitting Fixes:
 - Decrease Dense layer in classification from 64 to 16
 - Increased the dropout before classification
- Global vs Local Context Fixes:
 - Added Dilation rate of 2 to the final two convolutional layer
 - Changed padding of first convolutional layer to "valid" instead of "same"



| Model Description | Validation Accuracy | Public Kaggle Accuracy | Private Kaggle Accuracy |
|-------------------|---------------------|------------------------|-------------------------|
| Model Baseline | ~0.85 | 0.780 | 0.773 |
| Version 2 | ~0.85 | 0.721 | 0.699 |

Version 2 Observations


- Test Results: 0.721
- Observations:
 - Edge artifacts on visualization heat maps
 - Suggests "valid" isn't an ideal choice here
 - Focus on craters still
 - Focus on small local contexts still
- Significant changes need to occur for the third version

Prob: 0.04084106


Original



Grad-CAM



Overlay




Prob: 0.99883491


Original



Grad-CAM




Overlay




Prob: 0.99979275


Original



Grad-CAM




Overlay



Model and Hyperparameter Tuning Version 3

- Global vs Local Context Fixes
 - Increased kernel size from 3 to 5 of first convolutional layer
 - Added LeakyReLU as activation instead of ReLU. Allows negative values to propagate, although at a decreased rate
 - Average Pooling Instead of Max Pooling
 - Max Pooling was hiding the darker cracks
- Overfitting
 - Added regularization to the convolutional layers
 - Kept the decreased Dense Layer size
- Edge Artifacts
 - Change back to “same” padding




| Model Description | Validation Accuracy | Public Kaggle Accuracy | Private Kaggle Accuracy |
|-------------------|---------------------|------------------------|-------------------------|
| Model Baseline | ~0.85 | 0.780 | 0.773 |
| Version 2 | ~0.85 | 0.721 | 0.699 |
| Version 3 | ~0.86 | 0.795 | 0.771 |


Version 3 Observations and Results

- Test Results: 0.795
- Observations
 - Smaller overfitting gap between train and validation metric
 - Regularization
 - Ignores non-relevant craters/features
 - Has global context on cracks
 - Kernel size increase
 - Better at identifying thin cracks
 - Average Pooling
- Overall, best model so far and doesn't highlight confounding factors, only cracks.
- Explainable model when GradCAM is applied


Prob: 0.99347937



Prob: 0.24076895




Prob: 0.97978514



Data Augmentation

- Augmentation Layers
 - Random Contrast
 - Random Flip
 - Random Rotation
 - Random Zoom
 - Random Translation
 - Gaussian Noise
- Test Result: 0.713
- Hypothesis
 - Augmentation was too aggressive and cause the cracks to potentially be removed if they were near the edge.
 - Future work should tune this portion of the model



| Model Description | Validation Accuracy | Public Kaggle Accuracy | Private Kaggle Accuracy |
|---------------------|---------------------|------------------------|-------------------------|
| Model Baseline | ~0.85 | 0.780 | 0.773 |
| Version 2 | ~0.85 | 0.721 | 0.699 |
| Version 3 | ~0.86 | 0.795 | 0.771 |
| Augmented Version 3 | ~0.77 | 0.713 | 0.704 |

Recommendations, Use Cases, and Next Steps

- Results

- Convolutional Neural Network for Binary Classification of images
- Explainable model with additional GradCAM visualization
- Model is capable of identifying the entire crack in the global context
- Model is capable of ignoring the non-relevant, but confounding factors of craters

- Benefits

- Improves the ability to automate defect identification
- Less expensive than having human inspectors
- Enables earlier identification of defects; thus, potentially mitigating costly repairs and replacements

- Next Steps

- Have **spatially labeled** images with either bounding boxes or individual pixels
- Further hyperparameter tuning with the architecture and classification head to increase **accuracy** and **explainability**
- Further work with **data augmentation** to decrease overfitting and reduce the amount of training data needed
- Explore transfer learning of commonly used models

Conclusion


- Binary Classification Convolutional Neural Networks are able to still highlight the important features in the original image that is important to the classification
- The model was relatively fast to train and inference on a commercially available GPU
- Future work:
 - Improved spatial labeling with bounding boxes instead of binary labeling for entire image
 - Data Augmentation to further reduce overfitting
 - Hyperparameter tuning to further improve performance

Prob: 0.99992788


Original



Grad-CAM



Overlay



Prob: 0.2528871

Original



Grad-CAM



Overlay

